

Heart Disease Health Indicators Dataset

Xing Yang (Bill) Lan

Logistic Regression & Lasso

“Having a cholesterol
is bad”



Gradient Boosted Random Forest Ensemble

“Cholesterol is an
Ordinally encoded
categorical variable”



Everything is categorical, even age

HeartDiseaseorAttack		HighBP		HighChol		CholCheck		BMI	
Data type	numerical	Data type	numerical	Data type	numerical	Data type	numerical	Data type	numerical
Count	253680	Count	253680	Count	253680	Count	253680	Count	253680
Distinct	2	Distinct	2	Distinct	2	Distinct	2	Distinct	84
Missing	0	Missing	0	Missing	0	Missing	0	Missing	0
Numerical column		Numerical column		Numerical column		Numerical column		Numerical column	

Plot type

Bar

Color

select

X axis


Age

Y axis

HeartDiseaseorAttack


The dataframe has 253680 rows and 22 columns. 10000 rows were used for the plot.

[Export to chart cell](#) or [export to code cell](#) for processing the whole dataset and advanced configurations.



Interaction Effects

“Fruit bad”?



Convert results to Categorical to select larger Positives

```
def check(vec):
    if type(vec[0]) not in (int, float, np.float64):
        return np.ravel(vec)
    return vec

# returns the accuracy between a model's predictions and Y_test 0-1
def accuracy(Y_pr, Y_ts):
    Y_pr, Y_ts = check(Y_pr), check(Y_ts)
    return sum([1 if pr == ts else 0 for pr,ts in zip(Y_pr, Y_ts)]) / len(Y_pr)

# returns confusion data
def confusion(Y_pr, Y_ts):
    Y_pr, Y_ts = check(Y_pr), check(Y_ts)
    TP, FP, TN, FN = 0, 0, 0, 0
    dd = {(1,1):TP, (1,0):FP, (0,1):FN, (0,0):TN}
    for pr, ts in zip(Y_pr, Y_ts):
        dd[(pr,ts)] += 1
    return {"TP": dd[(1,1)], "FP": dd[(1,0)], "FN": dd[(0,1)], "TN": dd[(0,0)]}

# an arbitrary function that allows you to use regressors
# Y_test is 0-1 but you can just use a percentile cutoff of model's predictions to convert linear
# predictions into binary
def convertToBinary(Y_pr, cutoffPercentile):
    cutoff = np.percentile(Y_pr, cutoffPercentile)
    return [1.0 if pr > cutoff else 0.0 for pr in Y_pr]

# Run a model. Literally any model that does not need to one-hot-encode
def runModel(model, __X, __Y, cutoffPercentile=85):
    X_train, X_test, Y_train, Y_test = train_test_split(__X, __Y, test_size=0.30, random_state=1)
```

Lasso, Guessing 0, and Logistic Regression



```
runModel(Lasso(), X, Y)
```

```
Accuracy: 0.905813097866078
```

```
MSE: 0.094186902133922
```

```
Confusion: {'TP': 0, 'FP': 0, 'FN': 7168, 'TN': 68936}
```

```
accuracy([0]*len(Y_test), Y_test)
```

```
0.905813097866078
```

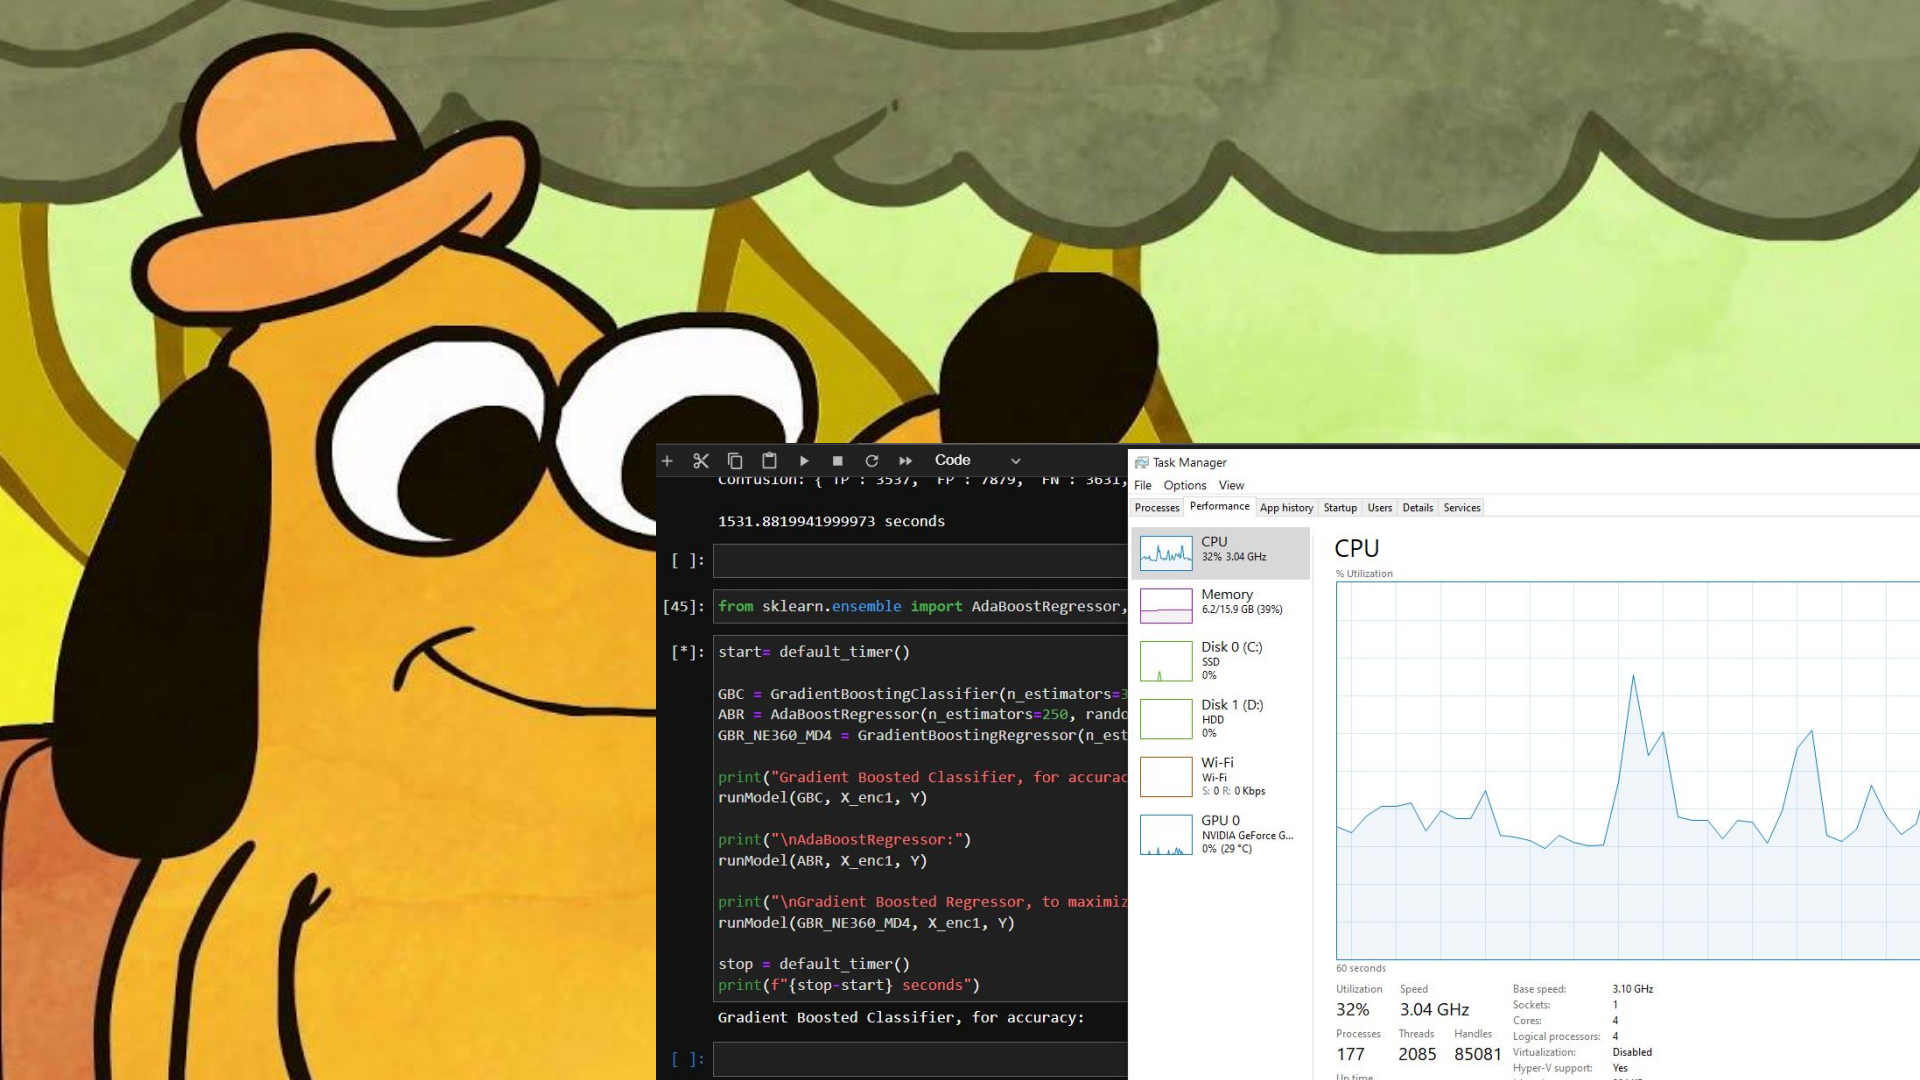
```
skLogiReg = LogisticRegression(max_iter = 2000)
```

```
runModel(skLogiReg, X, Y)
```

```
Accuracy: 0.9086118995059392
```

```
MSE: 0.09138810049406076
```

```
Confusion: {'TP': 916, 'FP': 703, 'FN': 6252, 'TN': 68233}
```

```
CONFUSION: { TP : 3537, FP : 7879, FN : 3631,

1531.8819941999973 seconds

[ ]:

[45]: from sklearn.ensemble import AdaBoostRegressor,

[*]: start = default_timer()

GBC = GradientBoostingClassifier(n_estimators=3
ABR = AdaBoostRegressor(n_estimators=250, rando
GBR_NE360_MD4 = GradientBoostingRegressor(n_est

print("Gradient Boosted Classifier, for accurac
runModel(GBC, X_enc1, Y)

print("\nAdaBoostRegressor:")
runModel(ABR, X_enc1, Y)

print("\nGradient Boosted Regressor, to maximiz
runModel(GBR_NE360_MD4, X_enc1, Y)

stop = default_timer()
print(f"{stop-start} seconds")


Gradient Boosted Classifier, for accuracy:


[ ]:
```


Task Manager


File Options View


Processes Performance App history Startup Users Details Services


 CPU
32% 3.04 GHz

 Memory
6.2/15.9 GB (39%)

 Disk 0 (C:) SSD
0%

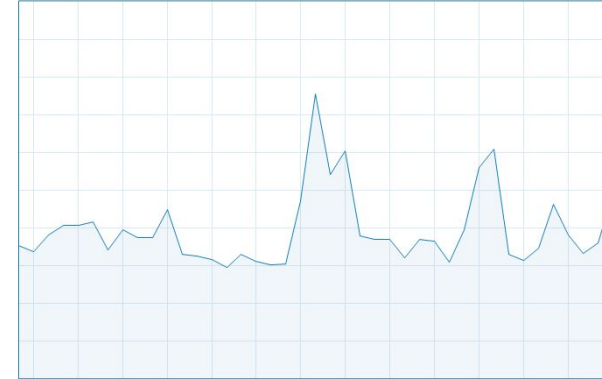
 Disk 1 (D:) HDD
0%

 Wi-Fi
Wi-Fi S: 0 R: 0 Kbps

 GPU 0
NVIDIA GeForce G... 0% (29 °C)

CPU

% Utilization



60 seconds

Utilization	Speed	Base speed:	3.10 GHz
32%	3.04 GHz	Sockets:	1
Processes	Threads	Cores:	4
177	2085	Logical processors:	4
Handles	85081	Virtualization:	Disabled
UIn Time		Hyper-V support:	Yes



```
Gradient Boosted Classifier, for accuracy:  
Accuracy: 0.9082702617470829  
MSE: 0.09172973825291707  
Confusion: {'TP': 823, 'FP': 636, 'FN': 6345, 'TN': 68300}
```



```
m12 = MLPClassifier(max_iter=400, random_state=1)  
runModel(m12, X_enc1, Y)  
  
Accuracy: 0.8895327446651949  
MSE: 0.110467255334805  
Confusion: {'TP': 1450, 'FP': 2689, 'FN': 5718, 'TN': 66247}
```



```
runModel(Lasso(alpha=0.01), X, Y, 90)  
  
Accuracy: 0.8725296962051929  
MSE: 0.12747030379480712  
Confusion: {'TP': 2539, 'FP': 5072, 'FN': 4629, 'TN': 63864}
```



Gradient Boosted Classifier, for accuracy:

Accuracy: 0.9082702617470829

MSE: 0.09172973825291707

Confusion: {'TP': 823, 'FP': 636, 'FN': 6345, 'TN': 68300}



Gradient Boosted Regressor, to maximize TP/FN:

Accuracy: 0.8572216966256702

MSE: 0.14277830337432987

Confusion: {'TP': 3859, 'FP': 7557, 'FN': 3309, 'TN': 61379}



Gradient Boosted Regressor at Cutoffs 85%, 90%, and 94%



GBR_NE460_MD3:

Accuracy: 0.8573793755912961

MSE: 0.14262062440870388

Confusion: {'TP': 3865, 'FP': 7551, 'FN': 3303, 'TN': 61385}

Accuracy: 0.8842504993167245

MSE: 0.11574950068327551

Confusion: {'TP': 2985, 'FP': 4626, 'FN': 4183, 'TN': 64310}

Accuracy: 0.9005177126038053

MSE: 0.09948228739619468

Confusion: {'TP': 2082, 'FP': 2485, 'FN': 5086, 'TN': 66451}

3431.6219225000023 seconds

